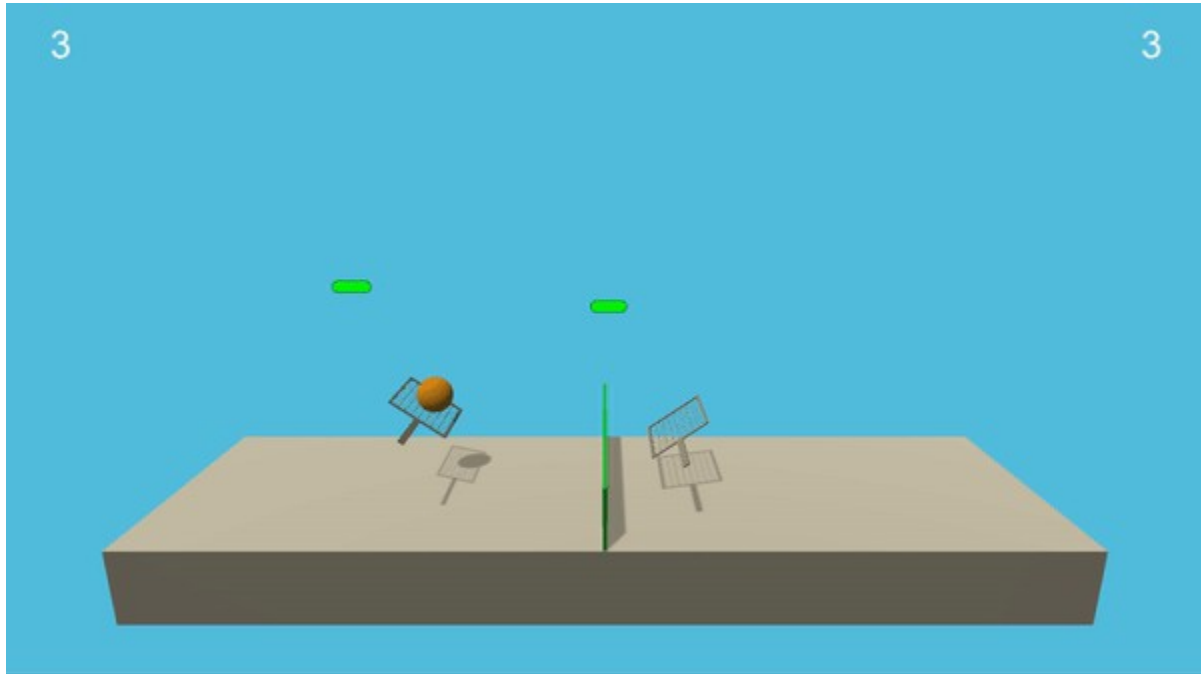


Deep Reinforcement Learning Nanodegree

Project 3- Collaboration-Competition.

REPORT



Introduction

This report provides a description of the implementation for the Deep Reinforcement Learning Nanodegree Project 3 to address the Multi-Agent Collaboration and Competition problem. In this environment, two agents control rackets to bounce a ball over a net. If an agent hits the ball over the net, it receives a reward of +0.1. If an agent lets a ball hit the ground or hits the ball out of bounds, it receives a reward of -0.01. Thus, the goal of each agent is to keep the ball in play.. Please refer to the [README.md](#) on this repository for more information regarding this environment and its installation.

Learning Algorithm

1. The Agent

After going through the Multi-agent lab to address the **Physical Deception** problem, I wanted to base this project of the code developed for this lab. Unfortunately, there were two main reasons I abandoned this idea:

- Udacity's Unity environment doesn't allow multiple environments (Threads). Unity crashes if you try to make a second instance of the environment...but on the Tennis environment on build by Unity ML agents they have solved this problem
- Also the way the **Physical Deception** Multi-Agent Deep Deterministic Policy Gradients (MADDPG) algorithm is structured assumes a single "brain" controlling 1 single agent. For that reason in the algorithm, the actions must be explicitly calculated for each agent separately. On the other side, on P2-Continuous Control (previous project) using 20 agents, the environment was using 1 single brain controlling all 20 agents. When I worked on this environment I already did the work needed for understanding and implementing the "multi-agent decentralized actor, centralized critic approach". That meant that the code from P2 needed minor changes to be able to work on P3
- After going through the **Physical Deception** problem lab code in detail, I knew that in reality, and knowing that I could not use multiple thread to accelerate the learning, there were essentially no major differences with the DDPG code developed on P2

Therefore I adapted the code from P2 to solve this problem

The Model Architecture Problem

Having already addressed my concerns about the architecture on P2 I decided to, first, try to run it with the same architecture since the problems were similar in their Observation and Actions Spaces. My initial trials showed success and therefore I decided to simply stick to my Model 2 architecture:

1. Model 2: with **3 hidden** layers (Actor) with a linear activation for each and Tanh activation for the output and with 2 **hidden** layers (Critic) with a Relu activation for each.

Recall that in my previous project I perform a set of tests trying different number of layers, different number of neurons per layer (filters on the Conv) and compared Relu and Leaky Relu and I noticed that the model with 3 hidden layers (originally found on Udacity's implementations) was learning faster and keeping a higher average of rewards

Summary of Additional Model Changes

ReLU Vs LeakyReLU: Leaky ReLU has a small slope for negative values, instead of altogether zero as ReLU does. The reason I wanted to try LeakyReLU was because it has two benefits over ReLU:

- It fixes the “dying ReLU” problem, as it doesn’t have zero-slope parts.
- It speeds up training. There is evidence that having the “mean activation” be close to 0 makes training faster. (It helps keep off-diagonal entries of the Fisher information matrix small, but you can safely ignore this.) Unlike ReLU, leaky ReLU is more “balanced,” and may therefore learn faster.

In this case, LeakyReLU did NOT show evidences of faster learning, although is possible to think that a different combination of my hyper-parameters could actually bring up LeakyReLU as a better choice.

Batch-normalization: The use of Batch normalization is a fascinating topic that started with me reading the paper that introduced it (<https://arxiv.org/abs/1502.03167>) Its benefits made clear sense after reading it but questions like: Should I apply Batch-Norm to just the input layer or to every input of every layer? Would it really make a difference? Researching on this topic was interesting and sometimes frustrating since I had no other choice but to perform a few tries...hoping that the differences observed were not going to be simply associated to the initial randomness of the weights. Batch-Normalization applied to every layer’s input showed the best performance. For details, please refer to model.py and model2.py files.

Gradient Clipping on the Critic Network: As it was clearly explained on the lectures, during ‘Training’ of a Deep Learning Model, we back-propagate our Gradients through the Network’s layers. Sometimes these gradients (tangent of the slopes) might be very large causing an overflow. Gradient clipping will ‘clip’ the gradients or cap them to a threshold value to prevent the gradients from getting too large. On the DDPG case and since the Critic Network uses the Actor’s best action approximation for training the Value/Critic Network, Clipping is only necessary on the Critic Net.

2. Experience Replay Vs Prioritized Experience Replay

The replay memory is a critical element on the DDPG (single and Multi-agent) implementation. In this case, using 2 agents, the replay memory was implemented “outside” the agent’s class allowing the agent’s experiences to be “memorized” together and later shared in the “Learning” process with the

other agents. This learning model was described as “multi-agent decentralized actor, centralized critic approach” on [“https://papers.nips.cc/paper/7217-multi-agent-actor-critic-for-mixed-cooperative-competitive-environments.pdf”](https://papers.nips.cc/paper/7217-multi-agent-actor-critic-for-mixed-cooperative-competitive-environments.pdf)

I wanted to test how Prioritized Experience Replay (PER) could improve my performance. It took me by surprise to see no improvement and even damaging the baseline performance. A quick look at this issue lead me to think that PER (<https://arxiv.org/abs/1511.05952>) uses the probability of a transition to occur, i.e. the number of times a transition occurs with respect the total number of transitions to calculate their “priority”. In a continuous action space, transitions can mathematically be infinite leading to priorities being totally normalized. At this point, when all memories (recorded transitions) have the same probability, sampling becomes exactly the same as in the basic Experience Replay, where we randomly select K samples. It seems only reasonable that prioritizing these memories using a different criteria of “importance” (more applicable to our case) might speed up the learning and even improve final performance. It make sense to think that using the “immediate” reward from every transaction might lead us to use and repeat the actions that lead us to a better reward. This is something I’d like to try in the future since time constraints to finish this assignment are not allowing me to do so now.

3. Other parameters of the MADDPG algorithm

Looking at “Class **args**” (defined on Tennis.ipynb) we can now focus and discuss how these arguments/parameters are used and how they affect the algorithm performance. Let me start by simply listing them below:

- seed = 777 Random seed
- disable_cuda = False Disable CUDA
- device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
- T_max = int(1e3) Number of training steps
- max_num_episodes = int(500) Max number of episodes
- hidden_1_size = 96 Network hidden layer 1 size
- hidden_2_size = 96 Network hidden layer 2 size
- hidden_3_size = 96 Network hidden layer 3 size
- noise_std = 0.05 Initial standard deviation of noise added to weights
- memory_capacity = int(1e6) Experience replay memory capacity
- batch_size = 256*1 Batch size: Number of memories will be sampled to learn
- learning_starts_ratio = 1/75 Number of steps before starting training => mem_capacity * ratio
- learning_frequency = 2 Steps before we sample from the replay Memory again
- priority_exponent = 0.5 Prioritized experience replay exponent (originally denoted α)

- `priority_weight = 0.4` Initial prioritized experience replay importance sampling weight
- `discount = 0.99` Discount factor
- `target_update = int(30)` # of steps after which to update target network
- `tau = 1e-3` Soft Update interpolation parameter
- `reward_clip = 1` Reward clipping (0 to disable)
- `lr_Actor = 1e-3` Learning rate - Actor
- `lr_Critic = 1e-3` Learning rate - Critic
- `adam_eps = 1e-08` Adam epsilon (Used for both Networks)
- `weight_decay = 0` Critic Optimizer Weight Decay

To keep the discussion short, and even though all parameters are relevant, I will only comment on those ones that I believe have a deeper impact on the final performance.

<code>T_max = int(1e3)</code>	# Number of training steps\n"
-------------------------------	-------------------------------

This parameter is definitely not that important, but I wanted to comment that not only is a good practice to limit the number of training steps taken on each episode to basically prevent it from getting into an infinite loop, but also it is interesting to see that in many "Toy Challenges" we need to use the "done" bit to be able to "reset" our environment in the case that "done" meant that we finished the task (failed or succeeded). So, for this reason I moved from: "while not np.any(dones):" to "while timestep<=T_Max:" in the Notebook code to make sure all episodes are consistently containing the same number of time steps/transitions.

<code>hidden_1_size = 96</code>	# Network hidden layer 1 size\n",
<code>hidden_2_size = 96</code>	# Network hidden layer 2 size\n",
<code>hidden_3_size = 96</code>	# Network hidden layer 3 size\n",

As mentioned before, these define the number of Neurons used in each of the 3 layers I finally used. 96 Neuron showed to be enough while 48 neurons were not enough to solve the task.

<code>memory_capacity = int(1e6)</code>	# Experience replay memory capacity\n",
<code>batch_size = 256*1</code>	# Batch size: Number of memories will be sampled to learn"
<code>learning_starts_ratio = 1/5</code>	# Number of steps before starting training = memory capacity * this ratio\n",
<code>learning_frequency = 20</code>	# Steps before we sample from the replay Memory again \n",

These are all Replay Memory/buffer related parameters. The "learning_starts_ratio" defines the number of transitions/steps that needed to be stored in memory before we start "learning". The "learning_frequency" defines the number of steps/transitions that we will have to go through (add to memory) after the condition to start sampling (learning) has been satisfied. In this way, we avoid sampling every timestep after this condition has been satisfied.

These parameters and their "proper" combination proved to be an important factor from the performance point of view. An art that took some time and testing to adjust. Please refer to P2 report for more details on how I tuned these parameters.

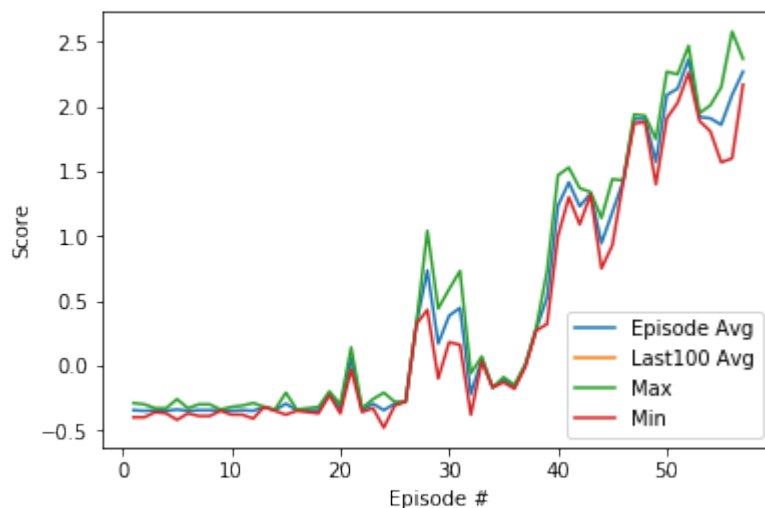
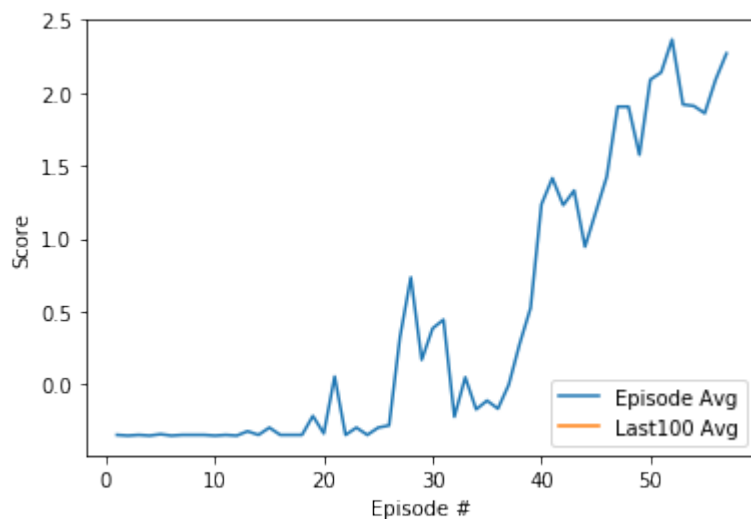
```
target_update = int(1)          # Number of steps after which to update
target network (Soft update for Actor & Critic)\n",
tau = 1e-3                      # Soft Update interpolation parameter\n",
```

These 2 parameters are related to when (as in how often) and how do we update the 4 Networks: The current and target networks for each the Actor and the Critic. In my last modifications I decided to make a soft update every time we passed all the conditions to “learn”

Results

On my first run and after just making sure my previous DDPG algorithm was adapted for this problem I got very successful results:

Episode 57 last 100 avg: 0.52 avg score: 2.27
 Solved in 57 episodes! last100_best_scores_average: 0.52, time
 taken(min): 19.38560005823771



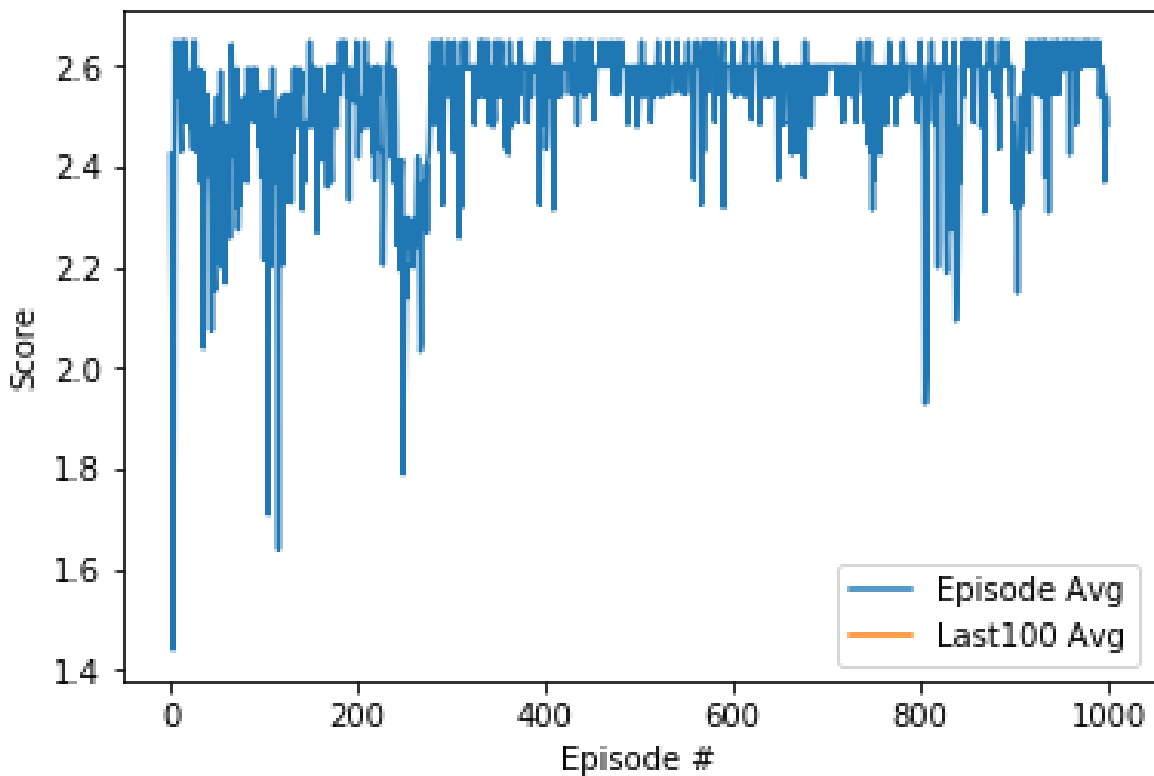
It took only 57 (constant length = 1000 steps) episodes to solve the challenge. That means that the algorithm needed 57,000 steps to solve the challenge


I decided not to stop the training when the challenge was solved (Average for the last 100 “best” agent episodes ≥ 0.5) and let the algorithm show me if it converges and has a plateau. Here are the results if let run for 1000 episodes at constant length of 1000 steps each episode, i.e 1,000,000 steps taken and used for learning: (Note: weights were pre-loaded from previous test that run 100 episodes)

Episode 10	last 100 avg: 2.40	avg score: 2.60
Episode 20	last 100 avg: 2.51	avg score: 2.54
Episode 30	last 100 avg: 2.53	avg score: 2.44
Episode 40	last 100 avg: 2.54	avg score: 2.54
Episode 50	last 100 avg: 2.52	avg score: 2.54
Episode 60	last 100 avg: 2.51	avg score: 2.49
Episode 70	last 100 avg: 2.51	avg score: 2.49
Episode 80	last 100 avg: 2.51	avg score: 2.39
Episode 90	last 100 avg: 2.52	avg score: 2.60
Episode 100	last 100 avg: 2.52	avg score: 2.22
Episode 110	last 100 avg: 2.52	avg score: 2.60
Episode 120	last 100 avg: 2.51	avg score: 2.37
Episode 130	last 100 avg: 2.50	avg score: 2.54
Episode 140	last 100 avg: 2.50	avg score: 2.59
Episode 150	last 100 avg: 2.51	avg score: 2.49
Episode 160	last 100 avg: 2.52	avg score: 2.54
Episode 170	last 100 avg: 2.52	avg score: 2.60
Episode 180	last 100 avg: 2.53	avg score: 2.55
Episode 190	last 100 avg: 2.53	avg score: 2.60
Episode 200	last 100 avg: 2.54	avg score: 2.60
Episode 210	last 100 avg: 2.55	avg score: 2.48
Episode 220	last 100 avg: 2.56	avg score: 2.60
Episode 230	last 100 avg: 2.57	avg score: 2.60
Episode 240	last 100 avg: 2.57	avg score: 2.59
Episode 250	last 100 avg: 2.55	avg score: 1.97
Episode 260	last 100 avg: 2.54	avg score: 2.30
Episode 270	last 100 avg: 2.51	avg score: 2.37
Episode 280	last 100 avg: 2.51	avg score: 2.65
Episode 290	last 100 avg: 2.51	avg score: 2.65
Episode 300	last 100 avg: 2.50	avg score: 2.54
Episode 310	last 100 avg: 2.50	avg score: 2.49
Episode 320	last 100 avg: 2.51	avg score: 2.60
Episode 330	last 100 avg: 2.52	avg score: 2.65
Episode 340	last 100 avg: 2.52	avg score: 2.55
Episode 350	last 100 avg: 2.54	avg score: 2.55
Episode 360	last 100 avg: 2.56	avg score: 2.49
Episode 370	last 100 avg: 2.59	avg score: 2.60
Episode 380	last 100 avg: 2.60	avg score: 2.60
Episode 390	last 100 avg: 2.60	avg score: 2.60

Episode 400	last 100 avg:	2.60	avg score:	2.60
Episode 410	last 100 avg:	2.60	avg score:	2.60
Episode 420	last 100 avg:	2.60	avg score:	2.60
Episode 430	last 100 avg:	2.60	avg score:	2.60
Episode 440	last 100 avg:	2.60	avg score:	2.60
Episode 450	last 100 avg:	2.60	avg score:	2.60
Episode 460	last 100 avg:	2.60	avg score:	2.60
Episode 470	last 100 avg:	2.60	avg score:	2.60
Episode 480	last 100 avg:	2.60	avg score:	2.60
Episode 490	last 100 avg:	2.60	avg score:	2.60
Episode 500	last 100 avg:	2.60	avg score:	2.55
Episode 510	last 100 avg:	2.61	avg score:	2.55
Episode 520	last 100 avg:	2.61	avg score:	2.65
Episode 530	last 100 avg:	2.61	avg score:	2.60
Episode 540	last 100 avg:	2.60	avg score:	2.55
Episode 550	last 100 avg:	2.60	avg score:	2.60
Episode 560	last 100 avg:	2.60	avg score:	2.60
Episode 570	last 100 avg:	2.60	avg score:	2.55
Episode 580	last 100 avg:	2.60	avg score:	2.60
Episode 590	last 100 avg:	2.60	avg score:	2.32
Episode 600	last 100 avg:	2.60	avg score:	2.60
Episode 610	last 100 avg:	2.60	avg score:	2.60
Episode 620	last 100 avg:	2.60	avg score:	2.50
Episode 630	last 100 avg:	2.60	avg score:	2.65
Episode 640	last 100 avg:	2.60	avg score:	2.60
Episode 650	last 100 avg:	2.60	avg score:	2.60
Episode 660	last 100 avg:	2.60	avg score:	2.60
Episode 670	last 100 avg:	2.60	avg score:	2.55
Episode 680	last 100 avg:	2.60	avg score:	2.49
Episode 690	last 100 avg:	2.59	avg score:	2.60
Episode 700	last 100 avg:	2.59	avg score:	2.60
Episode 710	last 100 avg:	2.59	avg score:	2.60
Episode 720	last 100 avg:	2.59	avg score:	2.60
Episode 730	last 100 avg:	2.59	avg score:	2.55
Episode 740	last 100 avg:	2.59	avg score:	2.60
Episode 750	last 100 avg:	2.59	avg score:	2.60
Episode 760	last 100 avg:	2.59	avg score:	2.60
Episode 770	last 100 avg:	2.59	avg score:	2.55
Episode 780	last 100 avg:	2.59	avg score:	2.55
Episode 790	last 100 avg:	2.60	avg score:	2.60
Episode 800	last 100 avg:	2.60	avg score:	2.65
Episode 810	last 100 avg:	2.58	avg score:	2.60
Episode 820	last 100 avg:	2.58	avg score:	2.55
Episode 830	last 100 avg:	2.58	avg score:	2.55
Episode 840	last 100 avg:	2.56	avg score:	2.19
Episode 850	last 100 avg:	2.56	avg score:	2.60
Episode 860	last 100 avg:	2.57	avg score:	2.65
Episode 870	last 100 avg:	2.57	avg score:	2.60
Episode 880	last 100 avg:	2.57	avg score:	2.60

Episode 890	last 100 avg: 2.57	avg score: 2.60
Episode 900	last 100 avg: 2.57	avg score: 2.43
Episode 910	last 100 avg: 2.57	avg score: 2.54
Episode 920	last 100 avg: 2.57	avg score: 2.55
Episode 930	last 100 avg: 2.58	avg score: 2.65
Episode 940	last 100 avg: 2.59	avg score: 2.60
Episode 950	last 100 avg: 2.60	avg score: 2.60
Episode 960	last 100 avg: 2.60	avg score: 2.65
Episode 970	last 100 avg: 2.60	avg score: 2.60
Episode 980	last 100 avg: 2.60	avg score: 2.65
Episode 990	last 100 avg: 2.60	avg score: 2.65
Episode 1000	last 100 avg: 2.60	avg score: 2.49





Even though the plots are not very clear due to the “clothing” of the amount of data plotted, we can still clearly see that Agent converges to a max of 2.60 rewards points very early in about 200 episodes (recall that the agent was pre-loaded with weights already from 100 episodes training)

This plateau is simply due to the pre-fixed time length of the game. After certain time the game is over, otherwise, the agents would highly likely play “forever” without dropping the ball.

Please refer to the video included in this repository (/Video.Tennis.mp4) for more graphical behavior of the agents in action.

Ideas for Future Work

- I truly wanted to explore a multi-thread approach.
- I definitely need to Implement other algorithms to:
 - o be able to compare performances.
 - o Understand better the true nature of multi-agent problems with other more complex “reward” functions that might show all; competitive, collaborative (like in this case) and mixed behaviors.
- Implement “Priority Replay Memory” with a “priority” criteria that fits this continuous actions and compare performances with normal (random sample) replay memory

- Investigate further the effects of other hyper parameters including: batch size, learning_starts_ratio and learning_frequency parameters described above.